Demo sprint 1

HPC scheduling with Kubernetes

Nidhi Shah, Yilin Xu, Siyuan Chen, Juhi Paliwal, Soufiane Jounaid

Agenda

- Background
- Problem and end goal
- Environment setup
- Use cases
- Burndown chart
- Plan for next sprint

Background

Kubernetes is a framework that runs distributed systems

- Pods: Containers containing application
- Nodes: Worker Machine that run Pods
- Control Plane: Manages Nodes

Kube-Scheduler selects the nodes that will run the new pods

Kubelet starts a pod

APIserver provides communication between components.

Limitations of HPC in Kubernetes:

- No Batch, Co-Scheduling
- No proper queuing system
- Kube-scheduler not adapted to HPC.

Kubernetes allows you to add your own scheduler.

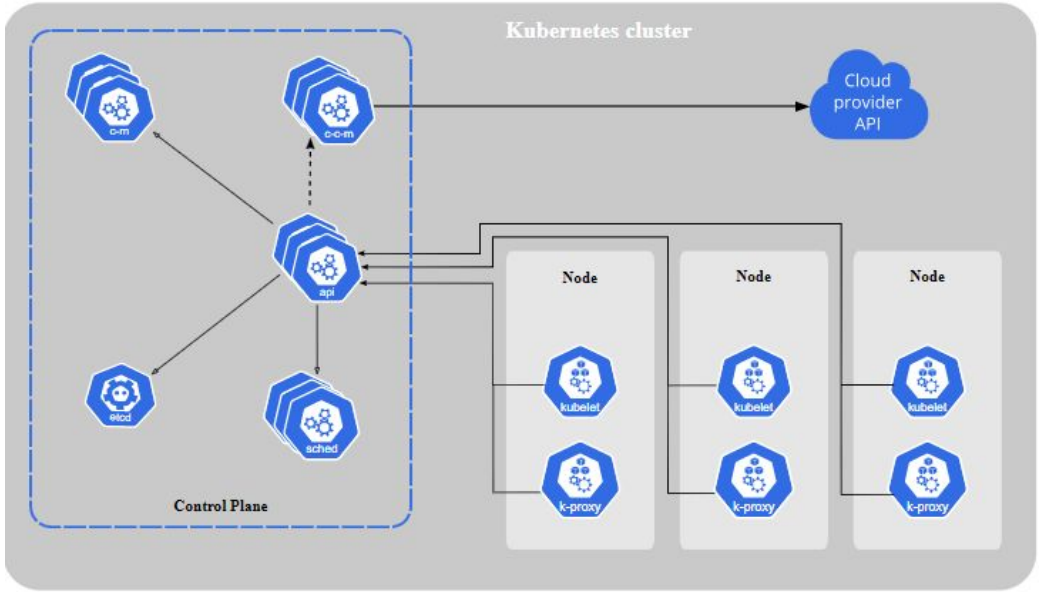

Components of Kubernetes

Kube-flux

Flux is a framework that uses a graph based scheduler called Fluxion.

Kube-Flux is an HPC scheduler that employs the Fluxion library to take scheduling decisions in Kubernetes.

Scheduling Cycle:

- Transform Kubernetes Pod specification to Flux Job Specification
- Submit placement request to Fluxion Library.
- The Fluxion library finds the node on the cluster and returns its name.
- Pass the selected node name to the API to bind with a pod.

Binding Cycle

- Pod are bound to a node

Problem

- Flux is not in charge of starting the actual pod, flux simply submits the pod to the Kubernetes APIserver which in turn works with Kublet to start the pod.
- Since Flux is not directly responsible for starting pods, it is not directly aware of their state on the cluster.
- If a pod scheduled by kube-flux finishes or is cancelled, kubernetes will remove said pod from the cluster and free the resources; this information however does not get propagated back to Kube-flux.
- Kube-Flux now sees a resource that has been freed as still occupied, which affects the correctness of future scheduling decisions.
- Consequently, since kubernetes does not receive state updates from the cluster, it is also unaware of the effects of scheduling decisions taken by other schedulers.

Illustration of the problem

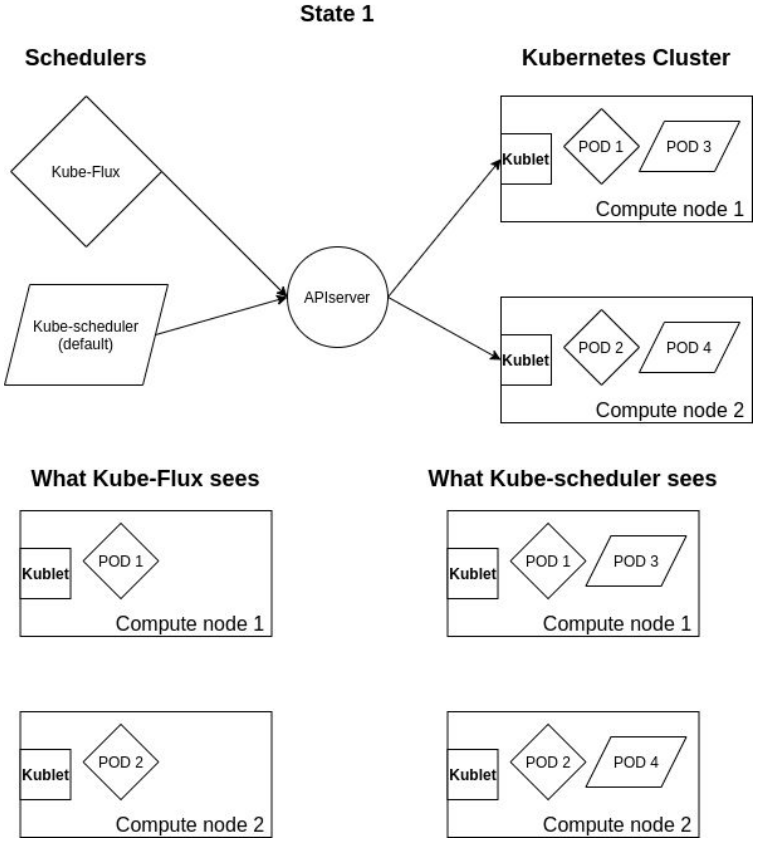

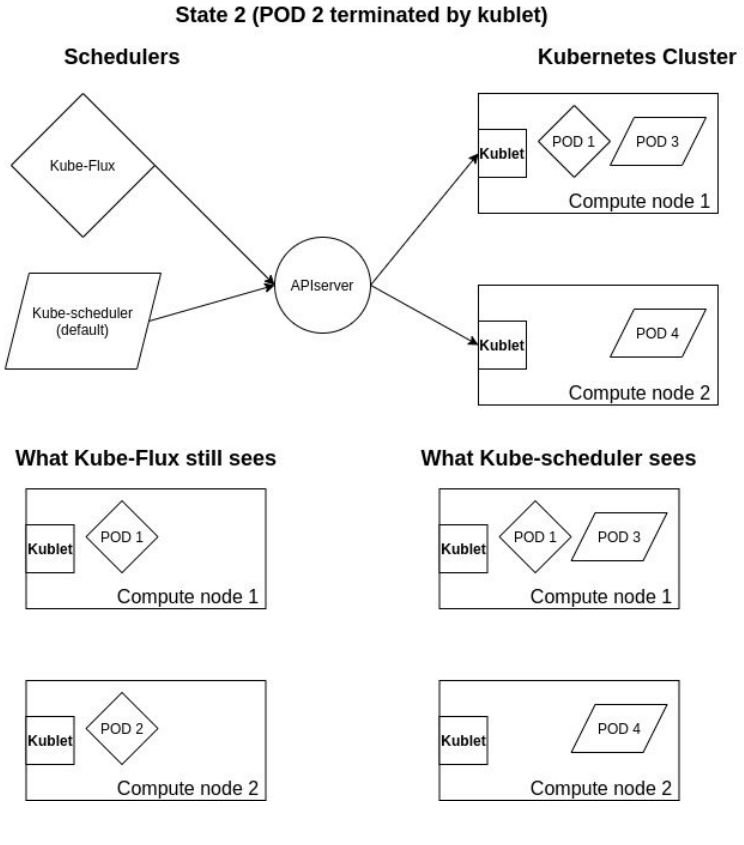

Goal of the project

- Full goal: Continuously synchronize the internal state of Flux with the state of the Kubernetes cluster of compute nodes.
    - Subgoal 1: Synchronize the Kube-flux scheduler to reflect changes in the state of its own scheduled pods.
    - Subgoal 2: Synchronize the Kube-flux scheduler to reflect changes in the state of pods scheduled by other schedulers.

Illustration of our goals (simplified)

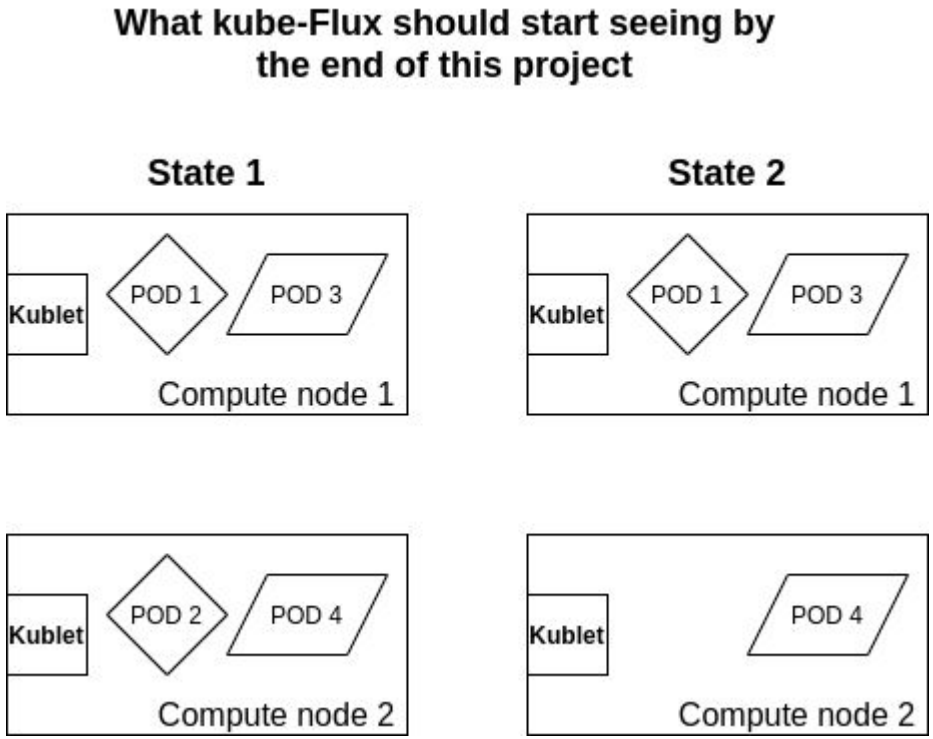

Potential solutions

- Obvious solution
  - Before each kube-flux scheduling round, we re-collect information about the compute nodes current utilization data and rebuild Kube-flux's internal graph using that information.
- Slightly less obvious solution
  - Collect the changes to about the state of the compute nodes and relay this information to Kube-flux which can then use the info to update its internal state without having to rebuild the graph.

Environment

- OS: Linux
- Dependencies:
  - Go: the language used in the development
  - Minikube: used to setup a local Kubernetes cluster
  - Helm: manage Kubernetes application
  - Docker: build images
  - CodeReady Containers (TBD): run local OpenShift clusters on a laptop
- Git Repository
  - Work together with mentors on a GitHub repo
- Local Repository
  - Clone the git repository
  - Test `make local-image`, the process takes about 8 mins

Environment

```
Setting up libboost-filesystem-dev:amd64 (1.71.0.0ubuntu4) ...
Setting up libboost-graph1.71-dev:amd64 (1.71.0-6ubuntu9) ...
Setting up libboost-graph-dev:amd64 (1.71.0.0ubuntu4) ...
Setting up libboost-system-dev:amd64 (1.71.0.0ubuntu4) ...
Setting up libczmq-dev:amd64 (4.2.0-2) ...
Processing triggers for libc-bin (2.32-0ubuntu3) ...
/sbin/ldconfig.real: /usr/local/lib/libflux-hostlist.so.1 is not a symbolic link

Reading package lists...
Building dependency tree...
Reading state information...
0 upgraded, 0 newly installed, 0 to remove and 0 not upgraded.
Removing intermediate container c95dc9897de0
 ---> 8fb513be335c
Step 48/51 : RUN mkdir -p /home/data/jobspecs/
 ---> Running in 446c31b66bbf
Removing intermediate container 446c31b66bbf
 ---> 16a47477b109
Step 49/51 : COPY flux-k8s/flux-plugin/manifests/kubeflux/sched-config.yaml /home/sched-config.yaml
 ---> 9e8dda23f3d5
Step 50/51 : WORKDIR /bin
 ---> Running in 56415043c551
Removing intermediate container 56415043c551
 ---> 32a9a7054ab8
Step 51/51 : CMD ["kube-scheduler"]
 ---> Running in c69aade6c0a6
Removing intermediate container c69aade6c0a6
 ---> e68758a16cc2
Successfully built e68758a16cc2
Successfully tagged localhost:5000/scheduler-plugins/kube-scheduler:latest
#docker build -f ./build/controller/Dockerfile --build-arg ARCH="amd64" -t localhost:5000/scheduler-plugins/controller:latest .
make local-image  0.61s user 0.65s system 0% cpu 7:51.50 total
→  scheduler-plugins git:(kubeflux) docker image ls
REPOSITORY                                          TAG       IMAGE ID       CREATED              SIZE
localhost:5000/scheduler-plugins/kube-scheduler     latest    e68758a16cc2   50 seconds ago       790MB
<none>                                              <none>    356b54915f49   About a minute ago   3.22GB
ubuntu                                              20.10     e508bd6d694e   2 months ago         79.4MB
→  scheduler-plugins git:(kubeflux)
```

Use cases

- Allow HPC applications that can run on a cloud cluster as well as they can run on an HPC cluster to be scheduled on a cloud cluster.
- Run both HPC and cloud workloads on a unified cluster.

Burndown Chart

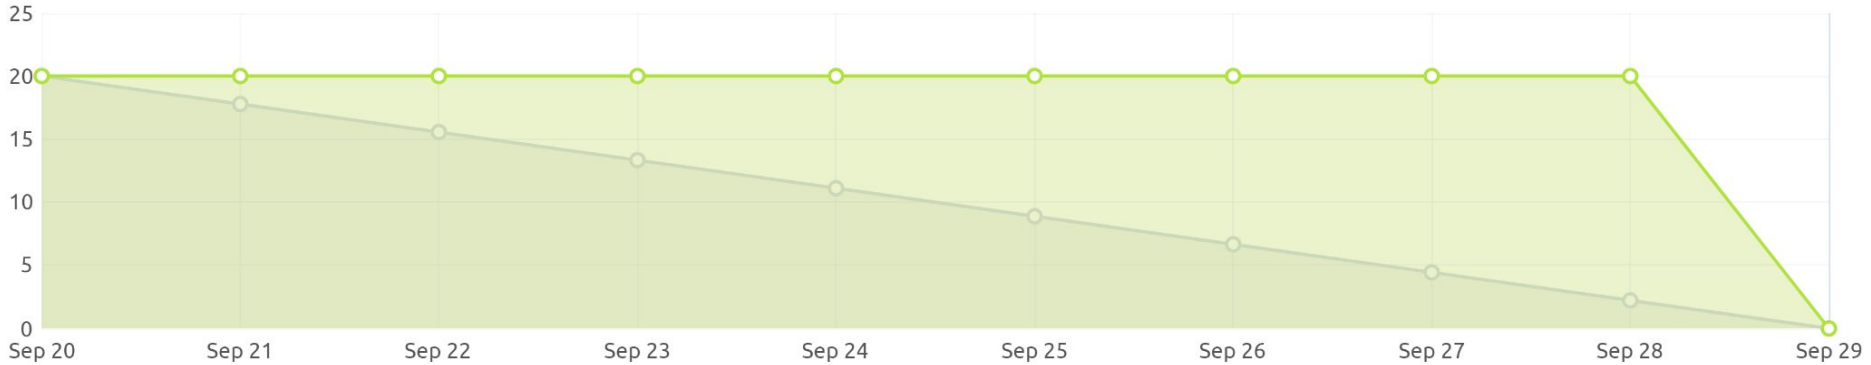

20 Points
8 Tasks

Sprint 2

- Setting up OpenShift resources
- Get familiar with Flux’s internal representation of the state of the Kubernetes cluster.
- Devise a testing plan by composing expected correct Flux states to test actual states against.

Thank You