

# HPC Scheduling with Kubernetes

Sprint 4

Siyuan Chen, Yilin Xu, Nidhi Shah, Soufiane Jounaid, Juhi Paliwal

Claudia Misale(IBM), Carlos Eduardo Arango Gutierrez (RedHat),  
Daniel Milroy (Lawrence Livermore National Laboratory)

# Content

- [Recap of the project](#)
- [Introduction to the state synchronization problem](#)
- [Potential solution: Operator](#)
- [Demo of the Operator](#)
- [Accomplished so far](#)
- [Burndown chart](#)
- [Plans for next sprint](#)

# About the Project so far

- State inconsistencies between the Kubernetes cluster and the Kube-Flux scheduler. This hinders utilization of the cluster, queued jobs wait forever.
  - Implemented an informer that carries updated cluster state information to the Kube-Flux resource graph for its own scheduled pods.
  - Exploring different solutions to the state synchronization (co-scheduling) issue for pods scheduled by multiple schedulers.

# Informer Solution

- KubeFlux schedules the pods on any available worker nodes. However it is not aware of any change in status of these allocated pods.
- This prohibits KubeFlux to request allocation of new pods on the same worker node.
- The implemented solution uses an informer that frequently watches the pods scheduled by KubeFlux and maintains the correct status of its pods.
- This solves the inconsistency of the state of resources at KubeFlux for its own scheduled pods.
- However, there still exists a problem with the pods scheduled by any other schedulers.

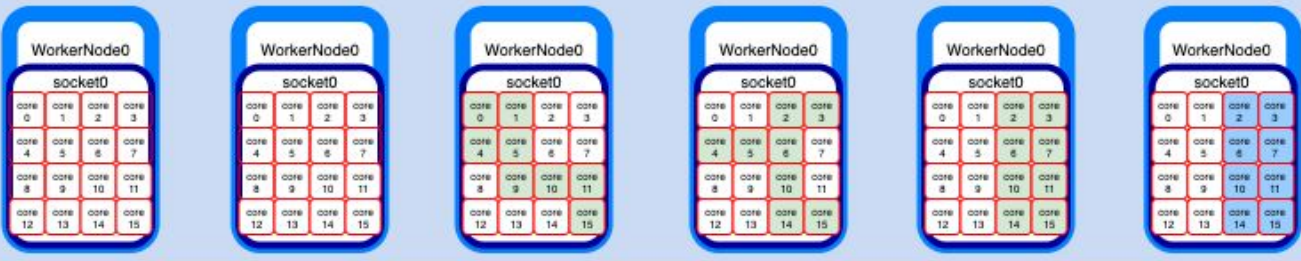
# Synchronization Problem across Schedulers



View from Kube Scheduler



View from KubeFlux



Legend:

- core i available
- core i In pending status
- core i In running status

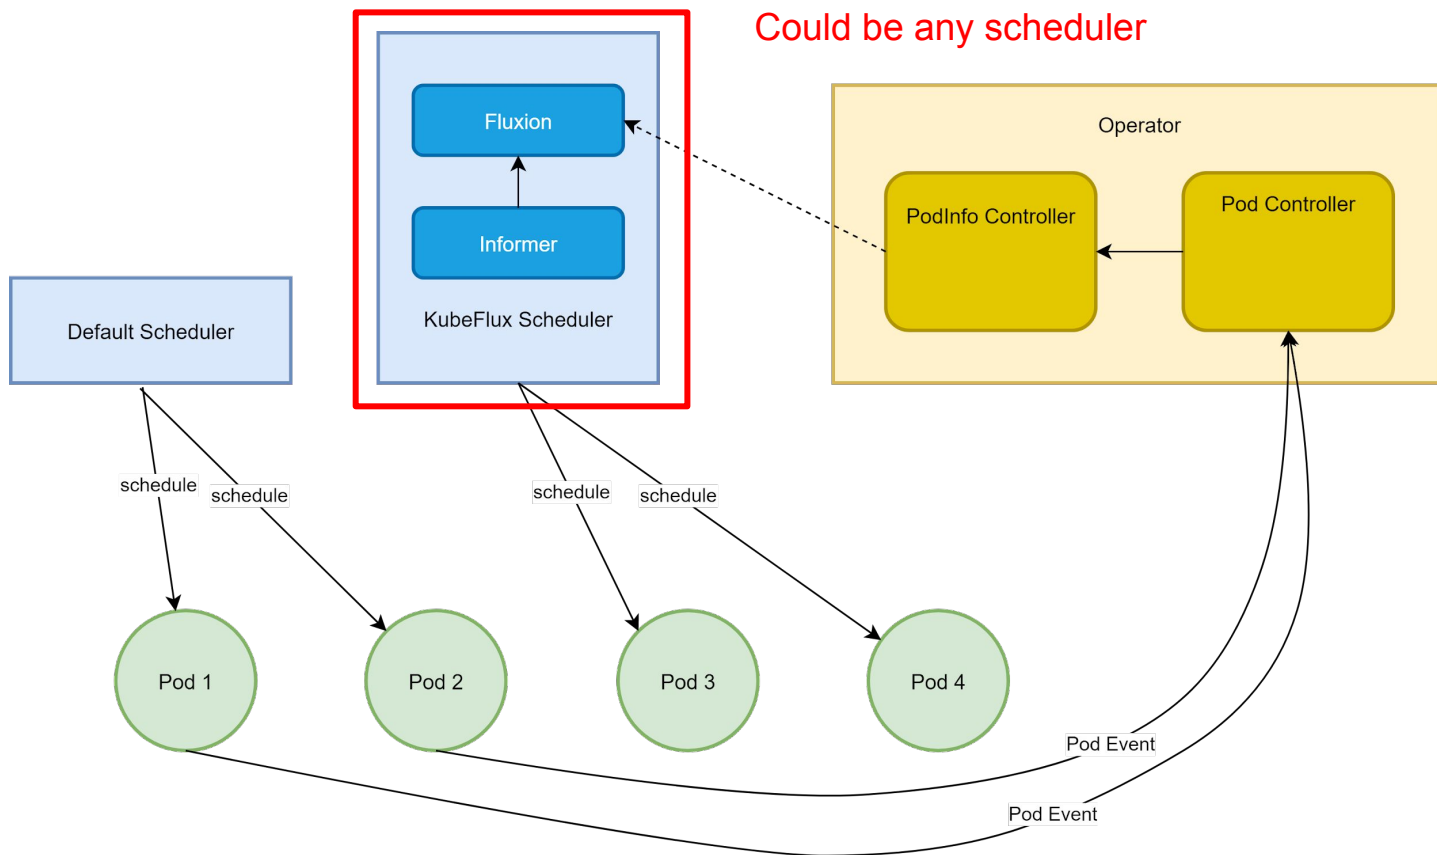
Cluster has 1 worker node with 1 cpu socket. It has 16 cores.

Kublet manages the allocation of pods for any scheduler and is aware of all the occupied resources.

# Main conceptual challenges

- Abstraction: Enabling co-scheduling of any two scheduler plugins.
  - The solution should offer a representation of how it records changes
  - Schedulers wanting to make use of co-scheduling to become aware of other schedulers can translate the changes recorded by the operator and update their own internal state accordingly.
  - Each component participating in co-scheduling is specific to a certain level. The interfaces have to be prescriptive of how to implement co-scheduling but not too restrictive.

# Potential Solution: Pod Controller



# Main technical challenges

- Kubernetes has 3 different ways of specifying pod resource needs.
  - Priority 1: Guaranteed Pod
  - Priority 2: Burstable Pod
  - Priority 3: Best-Effort Pod
- Resources usage fluctuations
- Imperative HPC management philosophy doesn't align with the declarative K8s management philosophy



# Demo Outline

1. Demo description
2. Deploy CRD and Operator
3. Run Pi Test with Default Scheduler
4. Check Generated Custom Resources


# Demo Description

- Environment: Cluster running on local machine
  - 2 nodes
    - 1 Master node: Kubernetes control plane
    - 1 worker node
      - 16 virtual cores in 1 socket and 1 memory node
- Tools:
  - Kind
  - Docker
- Task:
  - Deploy our operator and then deploy Pi calculation program
- Goal
  - Demonstrate Custom Resources are generated for pods running on a worker node

# Accomplished So Far

- Define CustomResourceDefinition (CRD)
- Test CRD with YAML file manually
- Create an operator through operator-sdk

# Burndown Chart



## Needs Info:

- Perform performance tests on a large cluster
- Analyze performance test results

## Plans for Next Sprint

- Demo informer in an open shift cluster (real world environment)
- Performance analysis of the Pod Informer